

Week 5: Strings and Files

CS211: Programming and Operating Systems

12 and 13 Feb, 2020



This week, in CS211:

- 1 Strings in C
- 2 `string.h`
- 3 String Output
- 4 String Input
- 5 Multidimensional Arrays
 - Arrays of Strings
- 6 Files
 - Getting started
 - Opening a file
 - Closing a file
 - An example
- 7 Reading from a file
- 8 Example 1: Using `fgets`
- 9 Example 2: Using `fgetc`

Strings in C

At the end of the last lecture, we studied *characters* in C. Now we will look at **strings**. Usually, these are thought of a collection of letters/characters that make up a word or a line of text.

The C language **does not actually have a string data type**. Instead, it uses arrays of type `char`.

If you make a declaration like:

```
char greeting[20]="Hello.  How are you?";
```

the system stores each character as an element of the array `greeting[]`.

Strings in C

Some Important Points:

1. In the above example we declared the array to be of length 20. Even though the string contains 19 characters, an extra ***string termination character*** `\0` (backslash zero) is added to show where the end of the string is.
2. Spaces or even new-line characters do not terminate a string. They are treated just like other characters.
3. Declarations are the only time we can use an “equals” to assign a value to a string. At all other times, we can modify the string one character at a time:
`greeting[0]='H'; greeting[1]='e'; ...`
4. Better still use `strcpy()` – the “string copy” function:
`strcpy(greeting, "Not too bad");`

Strings in C

The `strcpy()` is one of a collection of functions for dealing with strings. Its definition is to be found in the `string.h` header file. More of this later...

Example: *Write a function that determines the length of a string.*

Strings in C

00StrLength.c

```
#include <string.h>
int length(char *);

int main(void )
{
    char greeting[20];

    strcpy(greeting, "Hello. How are you?");

    printf("%s\n", greeting);

    printf("That message was %d chars long.\n",
           length(greeting) );

    return(0);
}
```

Strings in C

00StrLength.c

```
int length(char *str)
{
    int i, length=0;
    for (i=0; str[i] != '\0'; i++)
        length++;
    return(length);
}
```

string.h

Useful functions defined in `string.h` include:

strncpy

```
char *strncpy(char *dest, const char *source, int n);
```

Copies at most `n` character from the string in `source` to `dest`. The advantage is that we won't copy more characters to `dest` than is allowed

Example

```
char Code[6], Name[20]="Operating Systems";  
strcpy(Code, Name); // Bad! Unexpected Results  
strncpy(Code, Name, 6); // OK.
```


strcat

`strcat()`: Concatenate two strings, i.e., append one string onto the end of another. E.g,

```
char message1[30]="Hello.";
char message2[30]=" How are you?";
strcat(message1, message2);
```

Now `message1` contains "Hello. How are you?";

strcmp

`strcmp(char *s1, char *s2)`: **Comp**ares two stings. It returns an integer:

- 0 if they are the same,
- negative if s_1 is the first alphabetically
- positive if s_2 comes first

Example

```
char Name0[20], Name1[20], First[20];  
strncpy(Name0, "Richie", 20);\\  
strncpy(Name1, "Dennis", 20);\\  
  
if ( strcmp(Name0, Name1) > 0)  
    strncpy(First, Name1, 20);
```

strlen

`strlen` Takes a single (pointer to) a string as its argument and returns an integer equal to its `length` minus 1. (**Why -1?**).

strstr

```
char *strstr( char *haystack, char *needle);
```

Searches for the first occurrence of the string `needle` in `haystack`. It returns a pointer to the start of the matching substring. Moreover, if `needle` is **not** found in `haystack` it returns `NULL`.

Example:

```
if (strstr(Code, "CS") != NULL)
    printf("%s is a CS course\n", Code);
```

String Output

You all know how to use `printf()` with strings:

```
printf("%s%s\n", "Good morning ", name);
```

or

```
printf("%s%8s\n", "Good morning ", name);
```

In the second example the *field width* specifier is given. This causes the second string to be “padded” so that it takes up a total of 8 spaces. This is useful for tabulated output.

One could also use `puts()`: this prints the contents of a string followed by a new-line character.

String Input

Input is a more complicated issue, but there are three basic methods:

- `scanf("%s", name);` reads a the next “word” from the input buffer (usually the key board) and stores it in the array `name[]`. A word is a sequence of characters that does not include a space, tab or newline character.
- to get more control of the input, you could use `getchar()` within a loop:

```
printf("What is your name? ");  
for (i=0;  
      (myname[i] = getchar()) != '\n';  
      i++);  
myname[i]='\0';
```

String Input

- `gets(string)`: this reads a line a input and stores it all (except the `'\n'`) in the array pointed to by `string`. This would be very useful, except that `gets()` is known to be buggy and is best avoided.

From the Linux manual page from `gets()`:

BUGS

Never use `gets()`. Because it is impossible to tell without knowing the data in advance how many chars `gets()` will read, and because `gets()` will continue to store characters past the end of the buffer, it is extremely dangerous to use. It has been used to break computer security. Use `fgets()` instead.

String Input

- `fgets(string, n, stdin)`: reads in a line of text from the keyboard (standard input) and stores at most `n` characters in array `string`. The new line character is stored.

Which ever you use is a matter of choice. My preference is always to write functions that use `getchar()` and related functions, particularly if reading from a file.

Multidimensional Arrays

If an array (particularly of integers or floats) is like a mathematical vector, then how do we define a matrix?

A matrix is a two-dimensional array. For example, to declare a 3×4 matrix of floats, we would use the syntax:

```
float A[3][4];
```

So

$$A = \begin{pmatrix} A[0][0] & A[0][1] & A[0][2] & A[0][3] \\ A[1][0] & A[1][1] & A[1][2] & A[1][3] \\ A[2][0] & A[2][1] & A[2][2] & A[2][3] \end{pmatrix}$$

In general an $n \times m$ array is declared as

```
float A[n][m];
```

Multidimensional Arrays

If a program has the line:

```
int A[3][4];
```

What really happens is that the system creates **three** arrays, each of length **four**. More precisely, it

- declares 3 pointers to type `int`: `A[0]`, `A[1]`, and `A[2]`,
- space for storing an integer is allocated to each of the addresses `A[0]`, `A[0]+1`, `A[0]+2`, `A[0]+3`, `A[1]`, `A[1]+1`, ..., and `A[2]+3`.

.....
This means that if `A[] []` is declared as a two-dimensional 3×4 array, then the following are equivalent:

- `A[1][2]`
- `*(A[1] + 2)`
- `*(*(A + 1) + 2)`
- `*(&A[0][0] + 4 + 2)`

Multidimensional Arrays

01Matrix.c

```
#include <stdio.h>
6 int main(void )
{
    int A[3][4]={ {1,2,3,4}, {5,6,7,8}, {9,10,11,12}};

    printf("A[1][2] = %d\n", A[1][2]);
12 printf("(a[1]+2) = %d\n", *(A[1] + 2));
    printf("(*(A+1)+2) = %d\n", *( *(A + 1) + 2));
14 printf("(&A[0][0] + 4 + 2) = %d\n",
        *( &A[0][0] + 4 + 2));

    return(0);
18 }
```

Multidimensional Arrays

In another example , we'll sum all the entries of a 3×4 array.

02Sum_a_matrix.c

```
6 #include <stdio.h>
8 int sum(int A[][4]);
10 int main(void )
11 {
12     int n;
13     int A[3][4]={1,2,3,4,5,6,7,8,9,10,11,12};
14
15     n = sum(A);
16
17     printf("Sum of the entries in A is %d \n",n);
18     return(0);
19 }
```

Multidimensional Arrays

02Sum_a_matrix.c

```
22 int sum(int A[][4])  
23 {  
24     int i,j, ans=0;  
25     for (i=0; i < 3; i++)  
26         for (j=0; j< 4; j++)  
27             ans += A[i][j];  
28     return(ans);  
29 }
```

Important: Notice that this function is defined only for arrays of size 3×4 . Even if we passed *n* and *m* as arguments to the function, we would still have to declare that *a* has 4 columns.

Multidimensional arrays often occur when dealing with arrays of strings.

Recall that in C, a **string** (collection of characters) is stored as a **char** array.

```
char Name[20]="Ada Lovelace";
```

This means that we have declared **Names** to be an array of 15 characters:

- 'A' is stored in **Name[0]**
- 'd' is stored in **Name[1]**
- 'a' is stored in **Name[2]**
- ...
- 'c' is stored in **Name[10]**
- 'e' is stored in **Name[11]**
- and '\0' is stored in **Name[12]**.

The remaining components, **Name[13]**, ..., **Name[19]** are unused.

If a single string is stored as a character array, then an array of strings is an ***Array of Arrays of chars***, more often called a ***two dimensional array***.

Example

```
char Names[10][20];  
strcpy(Names[0], "A. Lovelace");  
strcpy(Names[1], "C. Babbage");  
...  
strcpy(Names[8], "D. Richie");  
strcpy(Names[9], "K. McNulty"); a
```

^aFor more about Donegal's greatest computer scientist, see https://en.wikipedia.org/wiki/Kathleen_Antonelli

We can think of this as a matrix, and visualise it as

	0	1	2	3	4	5	6	7	8	9	10	11
Name[0]	A	.		L	o	v	e	l	a	c	e	\0
Name[1]	C	.		B	a	b	b	a	g	e	\0	-
	⋮				⋮				⋮			⋮
Name[8]	D	.		R	i	c	h	i	e	\0	-	-
Name[9]	K	.		M	c	N	u	l	t	y	\0	-

.....

Clearly there is some waste of memory space. On another day, we might study the use of “ragged arrays” can avoid this.

Files

Most useful programs obtain their input from a **file**, and store their output to a file.

For example, in next week's lab (TBC) you'll develop a crossword helper that uses data stored in a file.

Further details can be found in Chap. 22 of King's "**C Programming**" or Chap 11 of Kelley and Pohl's "**A Book on C**".

Taking input from a file is not much different than taking input from the keyboard. All we do is:

- 1 Declare an identifier for the file, (`FILE *`)
- 2 open the file, (`fopen`)
- 3 read from it,
- 4 close the file. (`fclose`)

Declaring a ***File Identifier*** is easy:

```
FILE *datafile;
```

So `datafile` is now a pointer that we can associate with a file or, more generally, a `stream`.

```
fileptr = fopen(char *FileName, char *Mode);
```

The `fopen()` function is used for file opening. It takes two arguments: the `name` of the file to open and the ***mode*** it will operate in. A file pointer is returned.

The most important modes for file operation are `r` reading and `w` writing, but there is also `a` appending.

```
fileptr = fopen(char *FileName, char *Mode);
```

Read mode: "r"

Use `fopen(FileName, "r")` to open a file that we want to read from. It is assumed that the file already exists. If it doesn't, `NULL` is returned.

Example

```
FILE *infile;  
infile = fopen("OldFile.txt", "r");
```

Write mode: "w"

Use `fopen(FileName, "w")` to open a file we want to write to. If the file does **not** already exist, it is created. If it is already in the file system, the contents are deleted.

Example

```
FILE *outfile;  
outfile = fopen("NewFile.txt", "w");
```

There is also **append** mode: "a", used to to append data to end of the file. The file is opened in **write** mode, but new data is added to the end, i.e., its existing contents are not overwritten.

In our examples, we assume that

- The we only want to read from the file.
- That we know its name in advance.

So our code includes

```
FILE *fileptr;  
fileptr=fopen("list.txt", "r");
```

If the file can't be opened, `NULL` is returned.

When we are done, we should close the file

```
fclose(fileptr);
```

Example

Give a segment of code that prompts the user for name of an input file, and opens it in `r` read mode. If a file *cannot* be opened, an error should be returned.


```
int main( void)
{
    char infilename[20];
    FILE *infile;

    printf("Enter file to read from: ");
    scanf("%s", infilename);
    infile=fopen(infilename, "r");

    if (infile == NULL)
    {
        printf("Error: couldn't open for reading");
        return (EXIT_FAILURE);
    }
    else
        printf("Opened %s for reading\n", infilename);
```

Apart from `fopen` and `fclose`, the important functions for manipulating files are

- **Reading:** `fgetc` and `fgets` (also: `fscanf`)
- **Writing:** `fputc`, `fputs` and `fprintf`
- **Check and change file counter:** `rewind`, but also `ftell` and `fseek`.

Reading from a file

There are quite a number of functions for reading data from a file. We'll look at two functions: `fgetc()` and `fgets()`

- `fgets(string, n, fileptr)`

reads in a line of text from the *fileptr* stream.

and stores at most `n-1` characters in the array *string*.

Reading stops after an EOF or a newline. If a newline is read, it is stored in the string. A `\0` is stored after the last read character.

- `fgetc(fileptr)`

reads and returns the next character in the file. If the end of the file is reached, `EOF` is returned.

Also: `fscanf(fileptr, "%s", CharArray);`

works rather like `scanf()` except that the input stream is *fileptr* rather than *stdin*.

Example 1: Using `fgets`

Example 1: Write a function that counts the number of lines in a file using `fgets()`

Example 1: Using fgets

03CountLinesWithfgets.c

```
12 #include <stdio.h>
   #include <stdlib.h>
14 #include <string.h>

16 int file_length(FILE *);

18 int main( void)
   {
20     char FileName[30];
       FILE *file;

       strcpy(FileName, "03CountLinesWithfgets.c");
24     file=fopen(FileName, "r");

26     printf("%s has %d lines\n", FileName,
           file_length(file));
28     return(EXIT_SUCCESS);
   }
```

Example 1: Using fgets

03CountLinesWithfgets.c

```
26 int file_length(FILE *file)
   {
28     int lines;
       char dummy[100];

       rewind(file);

       lines=0;
34     while( fgets(dummy, 100, file) != NULL)
           lines++;

       rewind(file);

       return(lines);
40 }
```

Example 2: Using `fgetc`

We'll redo this example but using `fgetc`. It reads one character at a time so we'll just count the number of times a newline is read. Note that `EOF` — `End of File` — is returned when we try to read beyond the end of the file.

Example 2: Using fgetc

04CountLinesWithfgetc.c

```
26 int file_length(FILE *file)
   {
28     int lines;
       char c;

       rewind(file);

       lines=0;
34   do {
       c = fgetc(file);
36     if (c == '\n')
         lines++;
38   } while(c != EOF);

40   rewind(file);
       return(lines);
42 }
```