

Week 8: Sparse matrices in MATLAB

Table of Contents

1 "sparse" and "full"	1
2. Storage.....	2
3. Building sparse matrices.....	3
4. Are sparse matrices faster?.....	6

1 "sparse" and "full"

In MATLAB, the standard (non-sparse) format is called **"full"**.

To convert between sparse and full formats, use the `sparse()` and `full()` functions.

```
clear;
A = [1.0,    0,    0, -2.0,  0;
     0, 2.5, -1.0,    0,  0;
    -0.5,    0,    0,  12, -6;
     0,    0,    0,  1.0,  0;
     0,   -3,    0    0  0]
```

```
A = 5x5
    1.0000         0         0    -2.0000         0
         0    2.5000   -1.0000         0         0
   -0.5000         0         0   12.0000   -6.0000
         0         0         0    1.0000         0
         0   -3.0000         0         0         0
```

```
B = sparse(A)
```

```
B =
(1,1)    1.0000
(3,1)   -0.5000
(2,2)    2.5000
(5,2)   -3.0000
(2,3)   -1.0000
(1,4)   -2.0000
(3,4)   12.0000
(4,4)    1.0000
(3,5)   -6.0000
```

Note the triplet format that is used to express the output of B:

(i, j) value

Also notice the order in which elements are stored: first the entries in Column 1, then Column 2, etc. And indices are always increasing.

To convert from sparse to full, use the `full()` function.

```
C = full(B)
```

```
C = 5x5
```

```

1.0000    0    0   -2.0000    0
0    2.5000   -1.0000    0    0
-0.5000    0    0   12.0000   -6.0000
0    0    0    1.0000    0
0   -3.0000    0    0    0

```

```
isequal(A,C)
```


```
ans = logical
      1
```

However, it is rarely a good idea to build a sparse matrix by first constructing a "full" one, and then converting to sparse.

Two other functions that are very useful when dealing with sparse matrices:

- `spy()` visualises the sparsity pattern
- `nnz()` reports the number of non-zeros.

```
spy(B)
```



```
nnz(B)
```

```
ans = 9
```

2. Storage

We should be able to verify that storing a matrix in sparse format can save space. The number of bytes required is given by the "whos()" function

```
whos A B
```

Name	Size	Bytes	Class	Attributes
A	5x5	200	double	
B	5x5	192	double	sparse

Let's try to larger matrices to see if there is much of a saving. This one is deliberately not a good candidate for a sparse matrix: all the entries are 1.

```
N = 10;
A2 = ones(N);
B2 = sparse(A2);
whos A2 B2
```

Name	Size	Bytes	Class	Attributes
A2	10x10	800	double	
B2	10x10	1688	double	sparse

This shows that if the matrix is "dense", we should not use sparse format.

Here, "density" refers to the proportion of non-zeros in a matrix: it is $d = \text{nnz}(A) / \text{numel}(A)$:

```
fprintf("The density of A is d=%f", nnz(A)/numel(A))
```

```
The density of A is d=0.360000
```

```
fprintf("The density of A2 (the matrix of 1s) is d=%f", nnz(A2)/numel(A2))
```

```
The density of A2 (the matrix of 1s) is d=1.000000
```

This leads to two questions:

1. can we predict how much memory is required to store a sparse matrix?
2. for what value of d to we save memory by storing a matrix as sparse?

To answer Q1: Although MATLAB presents a sparse matrix in triplet form, it actually uses Compressed Column. Also, it uses 8 bytes for both indices and values. So an $N \times N$ matrix will need $2 \times 8 \times \text{NNZ} + 8(N + 1)$ bytes:

```
whos B
```

Name	Size	Bytes	Class	Attributes
B	5x5	192	double	sparse

```
2*8*nnz(B) + 8*(length(B)+1)
```


```
ans = 192
```

3. Building sparse matrices.

There are several functions that automatically return sparse matrices, including

- `speye(N)` : sparse $N \times N$ identity matrix
- `sprand(m,n,d)` : a random $m \times n$ matrix with specified density, d .

```
A3 = sprand(8,8, 0.5);  
spy(A3); title('sprand(8,8, 0.5)')
```




However, usually we construct matrices in a more deterministic way.

One way to do this is to make an empty sparse matrix. Then, when we add elements to it, it stays sparse.

In the following example, we'll make a tridiagonal matrix.

```
N = 5;  
A4 = sparse(N);  
A4(1,1)=1;  
for i=2:N-1  
    A4(i,i-1:i+1) = [-1,2,-1];  
end  
A4(N,[N-1,N])=[-1,2];  
spy(A4)
```



An alternative, and better, use of `sparse()` is to use triplet format:

```
A = sparse(I, J, X, N, N)
```

sets A to be the $N \times N$ matrix with $a_{I_k, J_k} = X_k$

Example:

```
I = 1:4
```

```
I = 1x4
    1    2    3    4
```

```
J = 2:5
```

```
J = 1x4
    2    3    4    5
```

```
X = [-1, 3, -2, 7]
```

```
X = 1x4
   -1    3   -2    7
```

```
A5 = sparse(I, J, X, 5, 5)
```

```
A5 =
    (1,2)    -1
    (2,3)     3
    (3,4)    -2
    (4,5)     7
```

```
full(A5)
```

```
ans = 5x5
      0     -1      0      0      0
      0      0      3      0      0
      0      0      0     -2      0
      0      0      0      0      7
      0      0      0      0      0
```

There are other related functions, including `spdiags()`, but I rarely find them to more useful than `sparse()`.

4. Are sparse matrices faster?

```
N = 2^8;
A_sparse = sparse(1:N, 1:N, 2) + ...
    sparse(1:N-1, 2:N, -1, N, N) + ...
    sparse(2:N, 1:N-1, -1, N, N);
%full(A_sparse)
%spy(A_sparse)
A_full = full(A_sparse);
b = ones(N,1);
tic; x = A_full\b; FullTime=toc;
tic; x = A_sparse\b; SparseTime=toc;
fprintf("N=%4d. Times: Sparse=%8.3gs, Full=%8.3gs (Speedup=%9.3g)\n", ...
    N, SparseTime, FullTime, FullTime/SparseTime);
```

```
N=2048. Times: Sparse=0.000118s, Full=    0.102s (Speedup=      861)
```